

iris

November 7, 2019

```
[43]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import statsmodels
```

```
[44]: auto = sns.load_dataset('mpg')
```

```
[45]: auto
# miles-per-gallon
```

```
[45]:      mpg  cylinders  displacement  horsepower  weight  acceleration  \
0    18.0          8         307.0         130.0   3504         12.0
1    15.0          8         350.0         165.0   3693         11.5
2    18.0          8         318.0         150.0   3436         11.0
3    16.0          8         304.0         150.0   3433         12.0
4    17.0          8         302.0         140.0   3449         10.5
..    ...          ...          ...          ...    ...          ...
393  27.0          4         140.0          86.0   2790         15.6
394  44.0          4          97.0          52.0   2130         24.6
395  32.0          4         135.0          84.0   2295         11.6
396  28.0          4         120.0          79.0   2625         18.6
397  31.0          4         119.0          82.0   2720         19.4
```

	model_year	origin	name
0	70	usa	chevrolet chevelle malibu
1	70	usa	buick skylark 320
2	70	usa	plymouth satellite
3	70	usa	amc rebel sst
4	70	usa	ford torino
..
393	82	usa	ford mustang gl
394	82	europe	vw pickup
395	82	usa	dodge rampage
396	82	usa	ford ranger
397	82	usa	chevy s-10

[398 rows x 9 columns]

```
[46]: type(auto)
```

```
[46]: pandas.core.frame.DataFrame
```

```
[47]: auto.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
mpg                398 non-null float64
cylinders          398 non-null int64
displacement       398 non-null float64
horsepower         392 non-null float64
weight             398 non-null int64
acceleration       398 non-null float64
model_year         398 non-null int64
origin             398 non-null object
name               398 non-null object
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB
```

```
[48]: auto['origin'].unique()
```

```
[48]: array(['usa', 'japan', 'europe'], dtype=object)
```

```
[49]: auto.max()
```

```
[49]: mpg                46.6
cylinders              8
displacement          455
horsepower            230
weight               5140
acceleration          24.8
model_year            82
origin                usa
name                 vw rabbit custom
dtype: object
```

```
[50]: auto.min()
```

```
[50]: mpg                9
cylinders              3
displacement          68
horsepower            46
weight               1613
acceleration           8
model_year            70
```

```
origin          europe
name            amc ambassador broughton
dtype: object
```

```
[51]: auto.mean()
```

```
[51]: mpg            23.514573
      cylinders      5.454774
      displacement  193.425879
      horsepower    104.469388
      weight        2970.424623
      acceleration   15.568090
      model_year     76.010050
      dtype: float64
```

```
[78]: auto.std()
```

```
[78]: mpg            7.815984
      cylinders      1.701004
      displacement  104.269838
      horsepower     38.491160
      weight        846.841774
      acceleration   2.757689
      model_year     3.697627
      dtype: float64
```

```
[83]: auto.columns
```

```
[83]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
          'acceleration', 'model_year', 'origin', 'name'],
          dtype='object')
```

```
[85]: list(auto.columns)
```

```
[85]: ['mpg',
      'cylinders',
      'displacement',
      'horsepower',
      'weight',
      'acceleration',
      'model_year',
      'origin',
      'name']
```

```
[86]: for col in auto.columns:
      print(col)
```

```
mpg
cylinders
displacement
horsepower
weight
acceleration
model_year
origin
name
```

```
[88]: for col in auto.columns:
      print(col)
      print(auto[col].std())
```

```
mpg
7.815984312565782
cylinders
1.7010042445332119
displacement
104.26983817119591
horsepower
38.49115993282849
weight
846.8417741973268
acceleration
2.757688929812676
model_year
3.697626646732623
origin
```

```
↳ -----
ValueError                                Traceback (most recent call↳
↳ last)

~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in f(values,↳
↳ axis, skipna, **kwds)
    119             else:
--> 120                 result = alt(values, axis=axis, skipna=skipna,↳
↳ **kwds)
    121             except Exception:
```

```
~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in↳
↳ nanvar(values, axis, skipna, ddof, mask)
```

```

767     # See https://en.wikipedia.org/wiki/
↳ Algorithms_for_calculating_variance
--> 768     avg = _ensure_numeric(values.sum(axis=axis, dtype=np.float64)) /
↳ count
769     if axis is not None:

~/.local/lib/python3.6/site-packages/numpy/core/_methods.py in _sum(a,
↳ axis, dtype, out, keepdims, initial, where)
37         initial=_NoValue, where=True):
---> 38     return umr_sum(a, axis, dtype, out, keepdims, initial, where)
39

```

ValueError: could not convert string to float: 'usa'

During handling of the above exception, another exception occurred:

```

ValueError                                Traceback (most recent call
↳ last)

~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in f(values,
↳ axis, skipna, **kwds)
122         try:
--> 123             result = alt(values, axis=axis, skipna=skipna,
↳ **kwds)
124         except ValueError as e:

~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in
↳ nanvar(values, axis, skipna, ddof, mask)
767     # See https://en.wikipedia.org/wiki/
↳ Algorithms_for_calculating_variance
--> 768     avg = _ensure_numeric(values.sum(axis=axis, dtype=np.float64)) /
↳ count
769     if axis is not None:

~/.local/lib/python3.6/site-packages/numpy/core/_methods.py in _sum(a,
↳ axis, dtype, out, keepdims, initial, where)
37         initial=_NoValue, where=True):
---> 38     return umr_sum(a, axis, dtype, out, keepdims, initial, where)
39

```

```
ValueError: could not convert string to float: 'usa'
```

During handling of the above exception, another exception occurred:

```
TypeError                                Traceback (most recent call
↳last)

~/local/lib/python3.6/site-packages/pandas/core/nanops.py in f(values,
↳axis, skipna, **kwds)
    119                 else:
--> 120                 result = alt(values, axis=axis, skipna=skipna,
↳**kwds)
    121             except Exception:

~/local/lib/python3.6/site-packages/pandas/core/nanops.py in
↳nanstd(values, axis, skipna, ddof, mask)
    710         """
--> 711         result = np.sqrt(nanvar(values, axis=axis, skipna=skipna,
↳ddof=ddof, mask=mask))
    712         return _wrap_results(result, values.dtype)

~/local/lib/python3.6/site-packages/pandas/core/nanops.py in _f(*args,
↳**kwargs)
    69             with np.errstate(invalid="ignore"):
---> 70                 return f(*args, **kwargs)
    71             except ValueError as e:

~/local/lib/python3.6/site-packages/pandas/core/nanops.py in f(values,
↳axis, skipna, **kwds)
    130                 if is_object_dtype(values):
--> 131                     raise TypeError(e)
    132                 raise

TypeError: could not convert string to float: 'usa'
```

During handling of the above exception, another exception occurred:

```
ValueError                                Traceback (most recent call
↳last)
```

```

~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in f(values,
↳axis, skipna, **kwds)
    119             else:
--> 120                 result = alt(values, axis=axis, skipna=skipna,
↳**kwds)
    121             except Exception:

```

```

~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in
↳nanvar(values, axis, skipna, ddof, mask)
    767     # See https://en.wikipedia.org/wiki/
↳Algorithms_for_calculating_variance
--> 768     avg = _ensure_numeric(values.sum(axis=axis, dtype=np.float64)) /
↳count
    769     if axis is not None:

```

```

~/.local/lib/python3.6/site-packages/numpy/core/_methods.py in _sum(a,
↳axis, dtype, out, keepdims, initial, where)
    37         initial=_NoValue, where=True):
---> 38     return umr_sum(a, axis, dtype, out, keepdims, initial, where)
    39

```

ValueError: could not convert string to float: 'usa'

During handling of the above exception, another exception occurred:

```

ValueError                                Traceback (most recent call
↳last)
~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in f(values,
↳axis, skipna, **kwds)
    122             try:
--> 123                 result = alt(values, axis=axis, skipna=skipna,
↳**kwds)
    124             except ValueError as e:

~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in
↳nanvar(values, axis, skipna, ddof, mask)
    767     # See https://en.wikipedia.org/wiki/
↳Algorithms_for_calculating_variance

```

```

--> 768     avg = _ensure_numeric(values.sum(axis=axis, dtype=np.float64)) /
↳count
    769     if axis is not None:

~/local/lib/python3.6/site-packages/numpy/core/_methods.py in _sum(a,
↳axis, dtype, out, keepdims, initial, where)
    37         initial=_NoValue, where=True):
--> 38     return umr_sum(a, axis, dtype, out, keepdims, initial, where)
    39

```

ValueError: could not convert string to float: 'usa'

During handling of the above exception, another exception occurred:

```

TypeError                                Traceback (most recent call
↳last)

```

```

<ipython-input-88-85729dd86791> in <module>
    1 for col in auto.columns:
    2     print(col)
----> 3     print(auto[col].std())

```

```

~/local/lib/python3.6/site-packages/pandas/core/generic.py in
↳stat_func(self, axis, skipna, level, ddof, numeric_only, **kwargs)
   11638         )
   11639         return self._reduce(
> 11640             f, name, axis=axis, numeric_only=numeric_only,
↳skipna=skipna, ddof=ddof
   11641         )
   11642

```

```

~/local/lib/python3.6/site-packages/pandas/core/series.py in
↳_reduce(self, op, name, axis, skipna, numeric_only, filter_type, **kwds)
   4088         )
   4089         with np.errstate(all="ignore"):
-> 4090             return op(delegate, skipna=skipna, **kwds)
   4091
   4092         # TODO(EA) dispatch to Index

```



```

~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in _f(*args,
↳ **kwargs)
    68         try:
    69             with np.errstate(invalid="ignore"):
---> 70                 return f(*args, **kwargs)
    71         except ValueError as e:
    72             # we want to transform an object array

```

```

~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in f(values,
↳ axis, skipna, **kwds)
    121         except Exception:
    122             try:
--> 123                 result = alt(values, axis=axis, skipna=skipna,
↳ **kwds)
    124             except ValueError as e:
    125                 # we want to transform an object array

```

```

~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in
↳ nanstd(values, axis, skipna, ddof, mask)
    709         1.0
    710         """
--> 711         result = np.sqrt(nanvar(values, axis=axis, skipna=skipna,
↳ ddof=ddof, mask=mask))
    712         return _wrap_results(result, values.dtype)
    713

```

```

~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in _f(*args,
↳ **kwargs)
    68         try:
    69             with np.errstate(invalid="ignore"):
---> 70                 return f(*args, **kwargs)
    71         except ValueError as e:
    72             # we want to transform an object array

```

```

~/.local/lib/python3.6/site-packages/pandas/core/nanops.py in f(values,
↳ axis, skipna, **kwds)
    129
    130         if is_object_dtype(values):
--> 131             raise TypeError(e)
    132         raise
    133

```

TypeError: could not convert string to float: 'usa'

```
[92]: for col in auto.columns:
      print(col)
      print(auto[col].dtypes)
      if (auto[col].dtypes == 'int64' or auto[col].dtypes == 'float64'):
          print(auto[col].std())
      print()
```

```
mpg
float64
7.815984312565782

cylinders
int64
1.7010042445332119

displacement
float64
104.26983817119591

horsepower
float64
38.49115993282849

weight
int64
846.8417741973268

acceleration
float64
2.757688929812676

model_year
int64
3.697626646732623

origin
object

name
object
```

```
[53]: auto['horsepower'][0]
```


```
[53]: 130.0
```

```
[77]: auto.loc[0]
```

```
[77]: mpg            18  
      cylinders      8  
      displacement  307  
      horsepower    130  
      weight        3504  
      acceleration   12  
      model_year     70  
      origin         usa  
      name           chevrolet chevelle malibu  
      Name: 0, dtype: object
```


```
[55]: sns.relplot(x = 'horsepower',  
                y = 'weight',  
                data = auto)
```

```
[55]: <seaborn.axisgrid.FacetGrid at 0x7fd60b2df240>
```




```
[56]: sns.relplot(x = 'horsepower',  
                y = 'weight',  
                hue = 'origin',  
                data = auto)
```

```
[56]: <seaborn.axisgrid.FacetGrid at 0x7fd60bcf1048>
```




```
[57]: sns.relplot(x = 'horsepower',  
                y = 'weight',  
                hue = 'origin',  
                col = 'origin',  
                data = auto)
```

```
[57]: <seaborn.axisgrid.FacetGrid at 0x7fd60c2550f0>
```



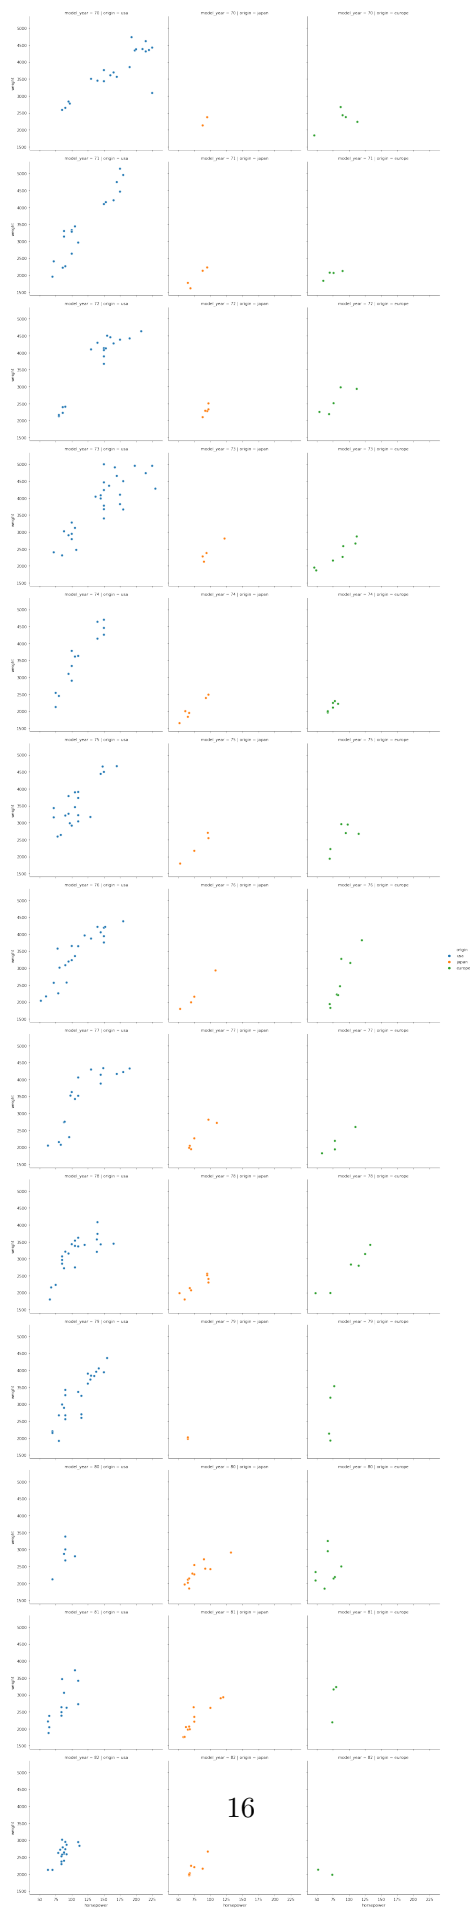
```
[58]: sns.relplot(x = 'horsepower',  
                y = 'weight',  
                hue = 'origin',  
                row = 'model_year',  
                data = auto)
```

```
[58]: <seaborn.axisgrid.FacetGrid at 0x7fd60be0a588>
```




```
[60]: sns.relplot(x = 'horsepower',  
                y = 'weight',  
                hue = 'origin',  
                row = 'model_year',  
                col = 'origin',  
                data = auto)
```

```
[60]: <seaborn.axisgrid.FacetGrid at 0x7fd60b4a7630>
```




```
[61]: sns.lmplot(x='weight', y='horsepower', data=auto)
```

```
[61]: <seaborn.axisgrid.FacetGrid at 0x7fd60ab1f748>
```



```
[62]: help(sns.lmplot)
```

Help on function lmplot in module seaborn.regression:

```
lmplot(x, y, data, hue=None, col=None, row=None, palette=None, col_wrap=None,
height=5, aspect=1, markers='o', sharex=True, sharey=True, hue_order=None,
col_order=None, row_order=None, legend=True, legend_out=True, x_estimator=None,
x_bins=None, x_ci='ci', scatter=True, fit_reg=True, ci=95, n_boot=1000,
units=None, order=1, logistic=False, lowess=False, robust=False, logx=False,
x_partial=None, y_partial=None, truncate=False, x_jitter=None, y_jitter=None,
scatter_kws=None, line_kws=None, size=None)
```

Plot data and regression model fits across a FacetGrid.

This function combines :func:`regplot` and :class:`FacetGrid`. It is intended as a convenient interface to fit regression models across conditional subsets of a dataset.

When thinking about how to assign variables to different facets, a general rule is that it makes sense to use ``hue`` for the most important comparison, followed by ``col`` and ``row``. However, always think about your particular dataset and the goals of the visualization you are creating.

There are a number of mutually exclusive options for estimating the regression model. See the :ref:`tutorial <regression_tutorial>` for more information.

The parameters to this function span most of the options in :class:`FacetGrid`, although there may be occasional cases where you will want to use that class and :func:`regplot` directly.

Parameters

x, y : strings, optional

Input variables; these should be column names in ``data``.

data : DataFrame

Tidy ("long-form") dataframe where each column is a variable and each row is an observation.

hue, col, row : strings

Variables that define subsets of the data, which will be drawn on separate facets in the grid. See the ``*_order`` parameters to control the order of levels of this variable.

palette : palette name, list, or dict, optional

Colors to use for the different levels of the ``hue`` variable. Should be something that can be interpreted by :func:`color_palette`, or a dictionary mapping hue levels to matplotlib colors.

col_wrap : int, optional

"Wrap" the column variable at this width, so that the column facets span multiple rows. Incompatible with a ``row`` facet.

height : scalar, optional

Height (in inches) of each facet. See also: ``aspect``.

aspect : scalar, optional

Aspect ratio of each facet, so that ``aspect * height`` gives the width of each facet in inches.

markers : matplotlib marker code or list of marker codes, optional

Markers for the scatterplot. If a list, each marker in the list will be used for each level of the ``hue`` variable.

share{x,y} : bool, 'col', or 'row' optional

If true, the facets will share y axes across columns and/or x axes across rows.

`{hue,col,row}_order` : lists, optional
 Order for the levels of the faceting variables. By default, this will be the order that the levels appear in ``data`` or, if the variables are pandas categoricals, the category order.

`legend` : bool, optional
 If ``True`` and there is a ``hue`` variable, add a legend.

`legend_out` : bool, optional
 If ``True``, the figure size will be extended, and the legend will be drawn outside the plot on the center right.

`x_estimator` : callable that maps vector -> scalar, optional
 Apply this function to each unique value of ``x`` and plot the resulting estimate. This is useful when ``x`` is a discrete variable. If ``x_ci`` is given, this estimate will be bootstrapped and a confidence interval will be drawn.

`x_bins` : int or vector, optional
 Bin the ``x`` variable into discrete bins and then estimate the central tendency and a confidence interval. This binning only influences how the scatterplot is drawn; the regression is still fit to the original data. This parameter is interpreted either as the number of evenly-sized (not necessary spaced) bins or the positions of the bin centers. When this parameter is used, it implies that the default of ``x_estimator`` is ``numpy.mean``.

`x_ci` : "ci", "sd", int in [0, 100] or None, optional
 Size of the confidence interval used when plotting a central tendency for discrete values of ``x``. If ``"ci"`` , defer to the value of the ``ci`` parameter. If ``"sd"`` , skip bootstrapping and show the standard deviation of the observations in each bin.

`scatter` : bool, optional
 If ``True``, draw a scatterplot with the underlying observations (or the ``x_estimator`` values).

`fit_reg` : bool, optional
 If ``True``, estimate and plot a regression model relating the ``x`` and ``y`` variables.

`ci` : int in [0, 100] or None, optional
 Size of the confidence interval for the regression estimate. This will be drawn using translucent bands around the regression line. The confidence interval is estimated using a bootstrap; for large datasets, it may be advisable to avoid that computation by setting this parameter to None.

`n_boot` : int, optional
 Number of bootstrap resamples used to estimate the ``ci``. The default value attempts to balance time and stability; you may want to increase this value for "final" versions of plots.

`units` : variable name in ``data``, optional
 If the ``x`` and ``y`` observations are nested within sampling units, those can be specified here. This will be taken into account when computing the confidence intervals by performing a multilevel bootstrap that resamples both units and observations (within unit). This does not

otherwise influence how the regression is estimated or drawn.

`order` : int, optional
 If `order` is greater than 1, use `numpy.polyfit` to estimate a polynomial regression.

`logistic` : bool, optional
 If `True`, assume that `y` is a binary variable and use `statsmodels` to estimate a logistic regression model. Note that this is substantially more computationally intensive than linear regression, so you may wish to decrease the number of bootstrap resamples (`n_boot`) or set `ci` to None.

`lowess` : bool, optional
 If `True`, use `statsmodels` to estimate a nonparametric lowess model (locally weighted linear regression). Note that confidence intervals cannot currently be drawn for this kind of model.

`robust` : bool, optional
 If `True`, use `statsmodels` to estimate a robust regression. This will de-weight outliers. Note that this is substantially more computationally intensive than standard linear regression, so you may wish to decrease the number of bootstrap resamples (`n_boot`) or set `ci` to None.

`logx` : bool, optional
 If `True`, estimate a linear regression of the form $y \sim \log(x)$, but plot the scatterplot and regression model in the input space. Note that `x` must be positive for this to work.

`{x,y}_partial` : strings in `data` or matrices
 Confounding variables to regress out of the `x` or `y` variables before plotting.

`truncate` : bool, optional
 By default, the regression line is drawn to fill the x axis limits after the scatterplot is drawn. If `truncate` is `True`, it will instead be bounded by the data limits.

`{x,y}_jitter` : floats, optional
 Add uniform random noise of this size to either the `x` or `y` variables. The noise is added to a copy of the data after fitting the regression, and only influences the look of the scatterplot. This can be helpful when plotting variables that take discrete values.

`{scatter,line}_kws` : dictionaries
 Additional keyword arguments to pass to `plt.scatter` and `plt.plot`.

See Also

`regplot` : Plot data and a conditional model fit.

`FacetGrid` : Subplot grid for plotting conditional relationships.

`pairplot` : Combine `:func:regplot` and `:class:PairGrid` (when used with `kind="reg"`).

Notes

The `:func:`regplot`` and `:func:`lmplot`` functions are closely related, but the former is an axes-level function while the latter is a figure-level function that combines `:func:`regplot`` and `:class:`FacetGrid``.

Examples

These examples focus on basic regression model plots to exhibit the various faceting options; see the `:func:`regplot`` docs for demonstrations of the other options for plotting the data and models. There are also other examples for how to manipulate plot using the returned object on the `:class:`FacetGrid`` docs.

Plot a simple linear relationship between two variables:

```
.. plot::
    :context: close-figs

    >>> import seaborn as sns; sns.set(color_codes=True)
    >>> tips = sns.load_dataset("tips")
    >>> g = sns.lmplot(x="total_bill", y="tip", data=tips)
```

Condition on a third variable and plot the levels in different colors:

```
.. plot::
    :context: close-figs

    >>> g = sns.lmplot(x="total_bill", y="tip", hue="smoker", data=tips)
```

Use different markers as well as colors so the plot will reproduce to black-and-white more easily:

```
.. plot::
    :context: close-figs

    >>> g = sns.lmplot(x="total_bill", y="tip", hue="smoker", data=tips,
    ...               markers=["o", "x"])
```

Use a different color palette:

```
.. plot::
    :context: close-figs

    >>> g = sns.lmplot(x="total_bill", y="tip", hue="smoker", data=tips,
    ...               palette="Set1")
```

Map ``hue`` levels to colors with a dictionary:

```
.. plot::
    :context: close-figs

    >>> g = sns.lmplot(x="total_bill", y="tip", hue="smoker", data=tips,
    ...                 palette=dict(Yes="g", No="m"))
```

Plot the levels of the third variable across different columns:

```
.. plot::
    :context: close-figs

    >>> g = sns.lmplot(x="total_bill", y="tip", col="smoker", data=tips)
```

Change the height and aspect ratio of the facets:

```
.. plot::
    :context: close-figs

    >>> g = sns.lmplot(x="size", y="total_bill", hue="day", col="day",
    ...                 data=tips, height=6, aspect=.4, x_jitter=.1)
```

Wrap the levels of the column variable into multiple rows:

```
.. plot::
    :context: close-figs

    >>> g = sns.lmplot(x="total_bill", y="tip", col="day", hue="day",
    ...                 data=tips, col_wrap=2, height=3)
```

Condition on two variables to make a full grid:

```
.. plot::
    :context: close-figs

    >>> g = sns.lmplot(x="total_bill", y="tip", row="sex", col="time",
    ...                 data=tips, height=3)
```

Use methods on the returned `:class:`FacetGrid`` instance to further tweak the plot:


```
.. plot::
    :context: close-figs

    >>> g = sns.lmplot(x="total_bill", y="tip", row="sex", col="time",
    ...                 data=tips, height=3)
    >>> g = (g.set_axis_labels("Total bill (US Dollars)", "Tip"))
```

```
...     .set(xlim=(0, 60), ylim=(0, 12),
...           xticks=[10, 30, 50], yticks=[2, 6, 10])
...     .fig.subplots_adjust(wspace=.02))
```


```
[63]: sns.lmplot(x = 'weight',
                 y = 'horsepower',
                 ci = None,
                 data = auto)
```

```
[63]: <seaborn.axisgrid.FacetGrid at 0x7fd60aa6cf28>
```




```
[64]: sns.lmplot(x = 'weight',
                 y = 'horsepower',
                 ci = 95,
                 data = auto)
```

```
[64]: <seaborn.axisgrid.FacetGrid at 0x7fd60aa33cc0>
```




```
[65]: sns.lmplot(x = 'weight',  
                y = 'horsepower',  
                hue = 'origin',  
                ci = None,  
                data = auto)
```

```
[65]: <seaborn.axisgrid.FacetGrid at 0x7fd60a9f1a90>
```


```
[66]: sns.lmplot(x = 'weight',  
               y = 'horsepower',  
               hue = 'origin',  
               ci = None,  
               col = 'origin',  
               data = auto)
```

```
[66]: <seaborn.axisgrid.FacetGrid at 0x7fd60a91c080>
```




```
[67]: sns.lmplot(x = 'weight',  
                y = 'horsepower',  
                ci = None,  
                col = 'model_year',  
                col_wrap = 4,  
                data = auto)
```

```
[67]: <seaborn.axisgrid.FacetGrid at 0x7fd60aa06550>
```




```
[68]: sns.lmplot(x = 'weight',
                 y = 'horsepower',
                 ci = None,
                 lowess = True,
                 data = auto)
```

```
[68]: <seaborn.axisgrid.FacetGrid at 0x7fd609c07e48>
```




```
[69]: sns.lmplot(x = 'weight',  
                y = 'horsepower',  
                hue = 'origin',  
                col = 'origin',  
                ci = None,  
                lowess = True,  
                data = auto)
```

```
[69]: <seaborn.axisgrid.FacetGrid at 0x7fd609cb2208>
```




```
[70]: plt.scatter(auto['weight'], auto['mpg'], label='data', color='red', marker='o',
        ↪s = 5)
```

```
[70]: <matplotlib.collections.PathCollection at 0x7fd609a5e1d0>
```




```
[71]: plt.scatter(auto['weight'], auto['mpg'], label='data', color='red', marker='o',
        ↪s = 5)
sns.regplot(x='weight', y='mpg', data=auto, color='blue', scatter=None,
        ↪label='order 1')
sns.regplot(x='weight', y='mpg', data=auto, color='green', order=2,
        ↪label='order 2', scatter=None)
plt.legend(loc='upper right')
```

[71]: <matplotlib.legend.Legend at 0x7fd609a76588>




```
[93]: sns.jointplot(x='horsepower',y='mpg',data=auto)
```

[93]: <seaborn.axisgrid.JointGrid at 0x7fd60b170ac8>




```
[95]: sns.jointplot(x='horsepower',y='mpg',kind = 'hex', data=auto)
```

```
[95]: <seaborn.axisgrid.JointGrid at 0x7fd60b45c0b8>
```



```
[96]: sns.jointplot(x='horsepower',y='mpg',kind = 'kde', data=auto)
```

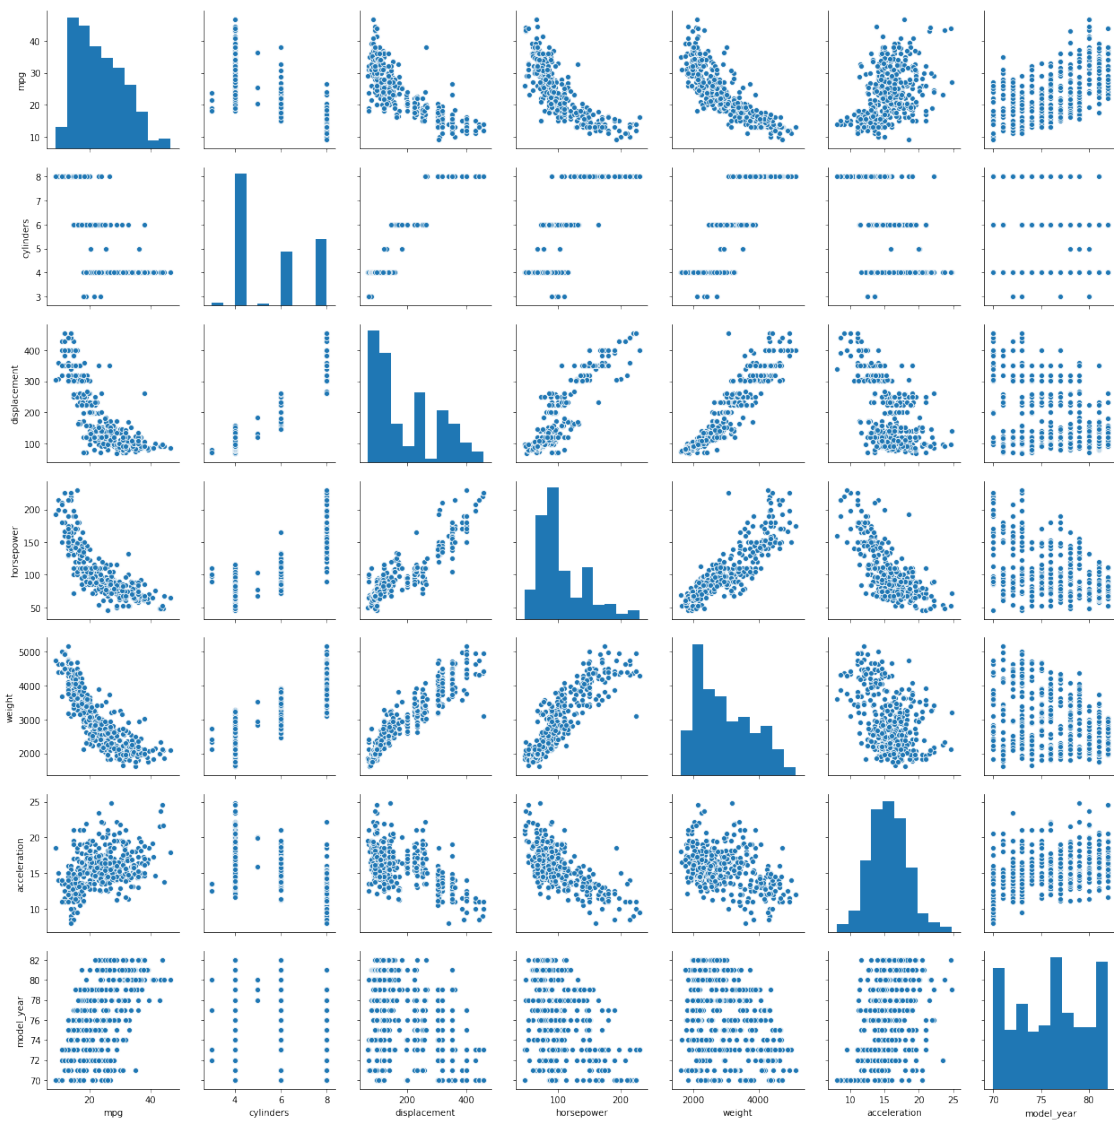
```
[96]: <seaborn.axisgrid.JointGrid at 0x7fd60b45c048>
```

```
[73]: sns.pairplot(data=auto)
```


```
/home/kroye/.local/lib/python3.6/site-packages/numpy/lib/histograms.py:829:  
RuntimeWarning: invalid value encountered in greater_equal  
    keep = (tmp_a >= first_edge)  
/home/kroye/.local/lib/python3.6/site-packages/numpy/lib/histograms.py:830:  
RuntimeWarning: invalid value encountered in less_equal  
    keep &= (tmp_a <= last_edge)
```

```
[73]: <seaborn.axisgrid.PairGrid at 0x7fd609789400>
```

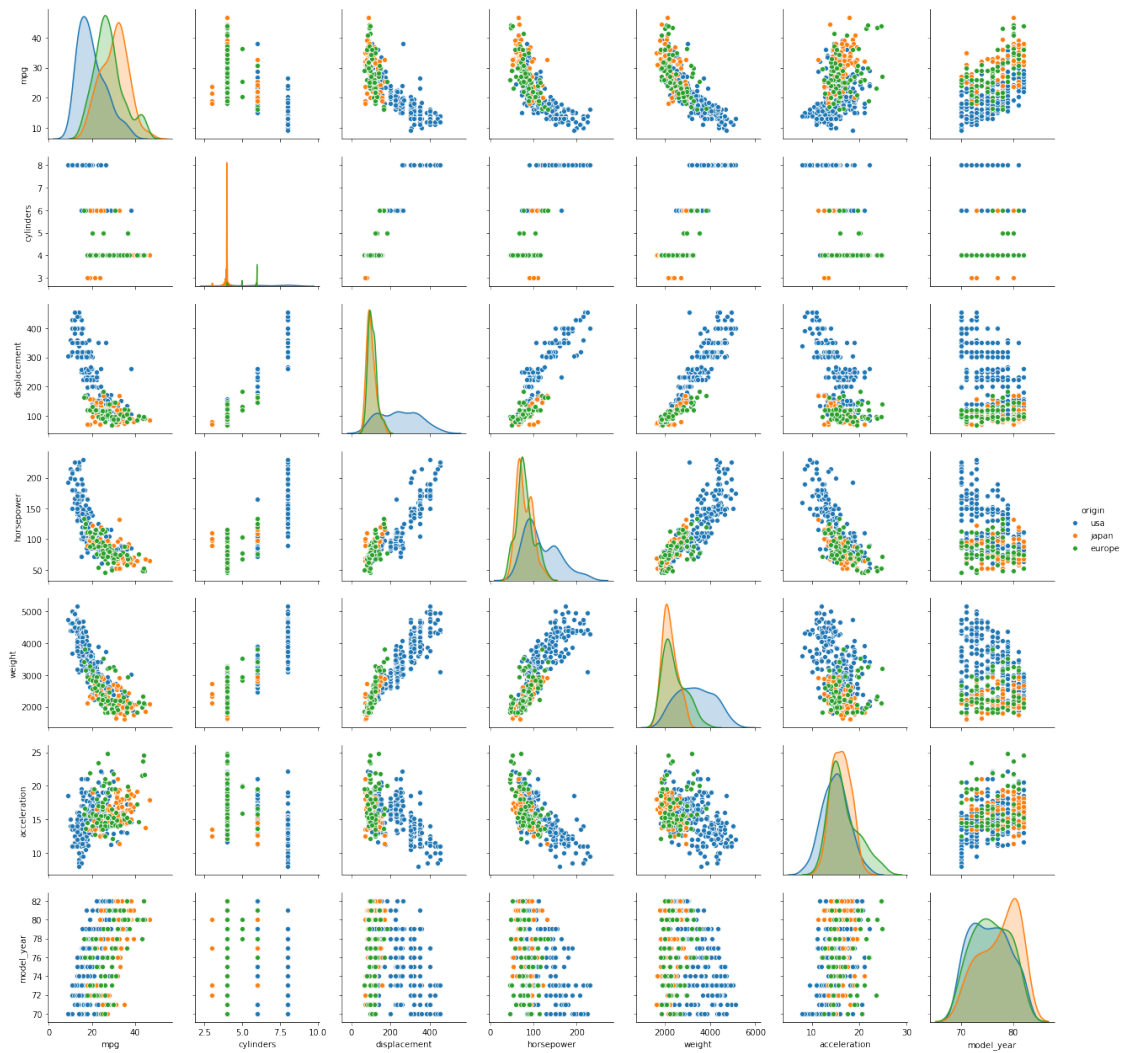


```
[74]: sns.pairplot(data=auto,hue='origin')
```

```
[74]: <seaborn.axisgrid.PairGrid at 0x7fd606c14400>
```




```
[75]: sns_fig = sns.pairplot(data=auto,hue='origin')
sns_fig.savefig('test.png')
```



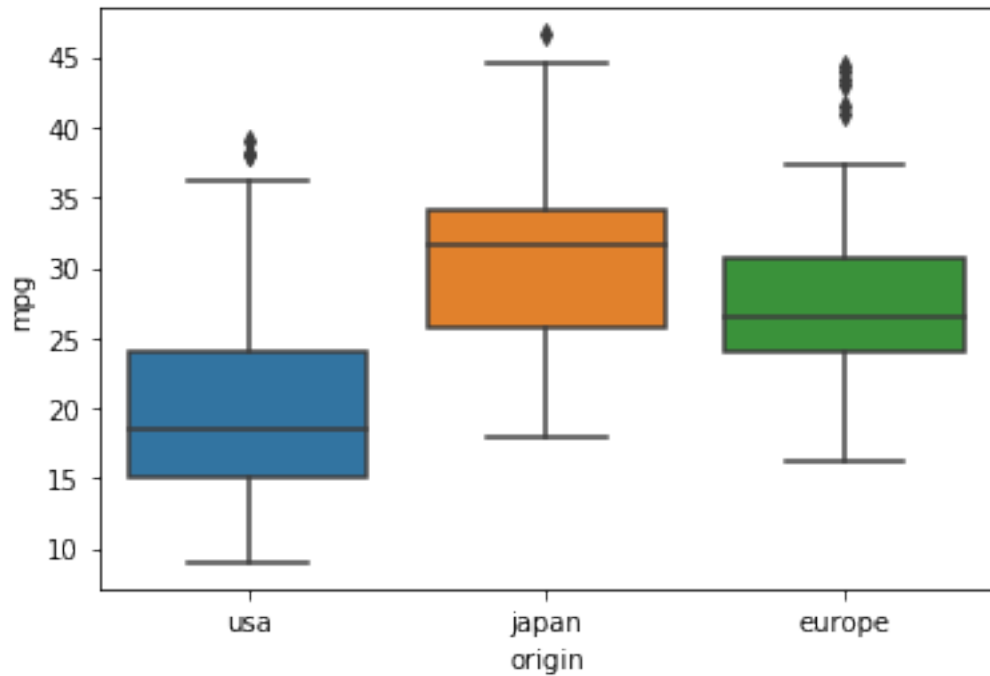
```
[76]: sns.relplot(x="horsepower", y="mpg", hue="origin", size="weight",
                sizes=(40, 400), alpha=.5, palette="muted",
                height=6, data=auto)
```

```
[76]: <seaborn.axisgrid.FacetGrid at 0x7fd6032bf6a0>
```





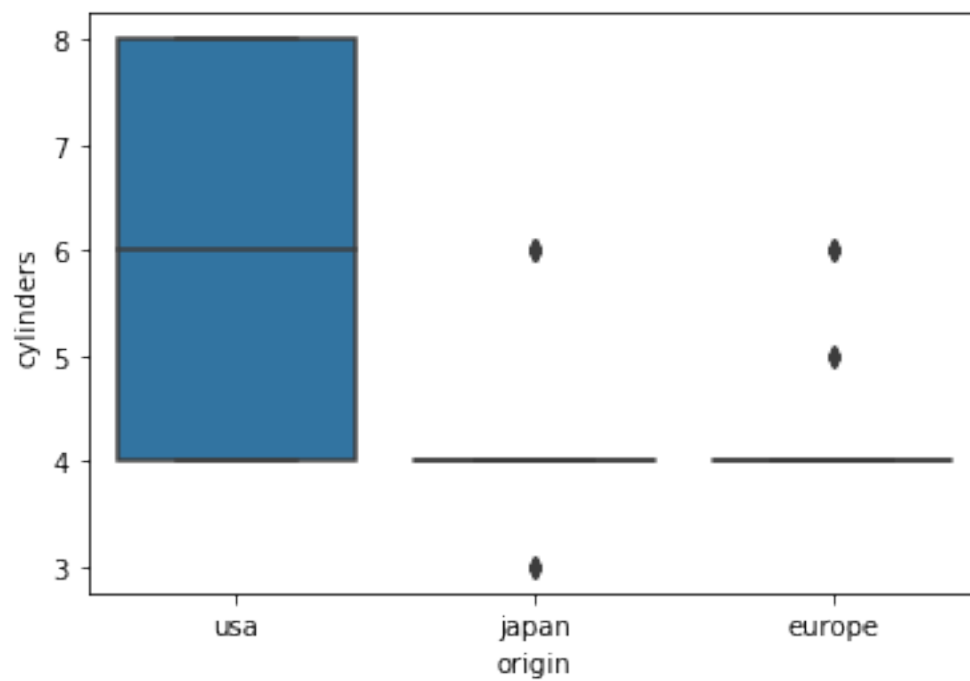
```
[100]: sns.boxplot(x="origin", y="mpg", data = auto)
```

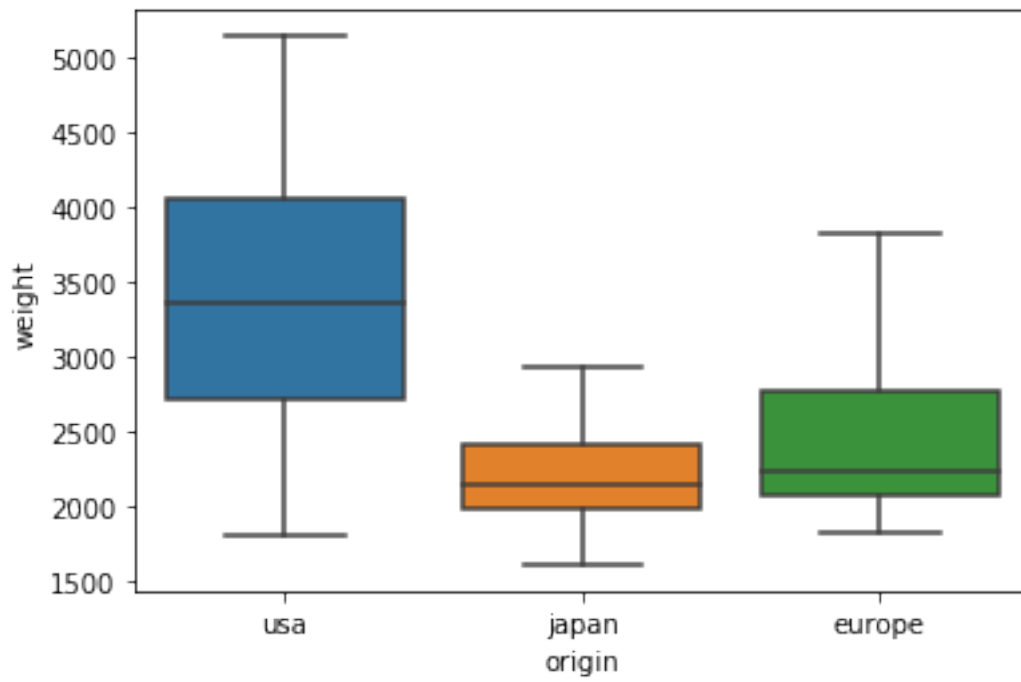

```
[100]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd60a685e48>
```





```
[106]: for col in auto.columns:
        if (auto[col].dtypes == 'int64' or auto[col].dtypes == 'float64'):
            sns.boxplot(x="origin", y = col, data = auto)
            plt.show()
```









[]: